
Session 3 (Database)

Mohamed Emary

January 9, 2025

We will work on training database with this schema:

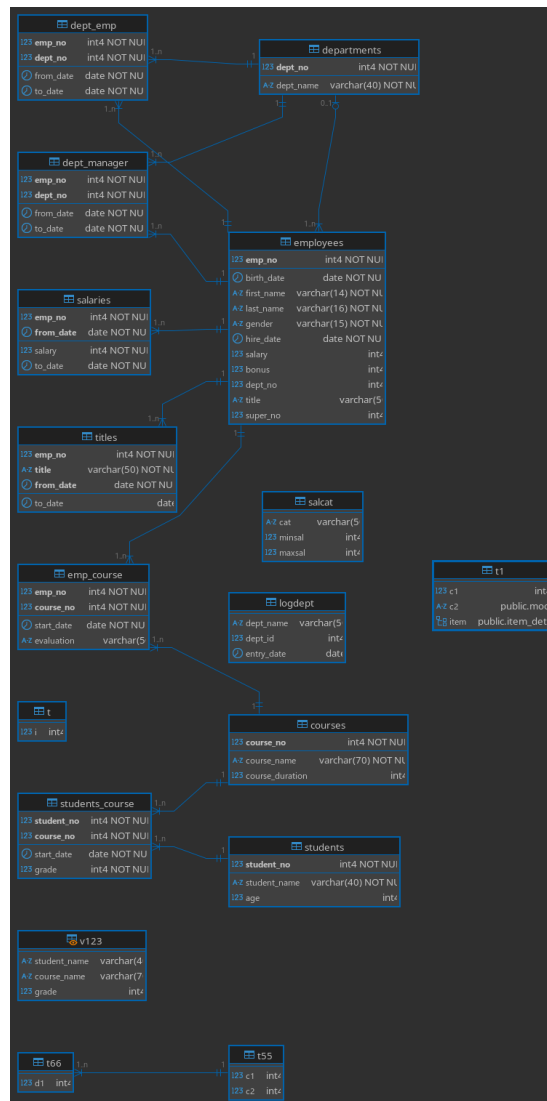


Figure 1: Database Schema

1 Database Join

1.1 Cross Join

Cross join returns all possible combinations of rows from two tables.

```
1 SELECT
2     first_name,
3     last_name,
4     salary,
5     dept_name
6 FROM
7     employees,
8     departments;
```

Cross join can be used via either `CROSS JOIN` or a comma , between table names as in the query above.

1.2 Inner Join

Inner join returns only the rows that have matching values in both tables.

For if we want to select values from both `employees` and `departments` tables where `dept_no` is the same in both tables, we can use the following query.

Since the relation between `employees` and `departments` tables is based on the `dept_no` column, we specified this column in the `WHERE` clause.

```
1 SELECT
2     first_name,
3     last_name,
4     salary,
5     dept_name
6 FROM
7     employees e,
8     departments d
9 WHERE
10    e.dept_no = d.dept_no;
```

The query above will return 997 rows, although the `employees` table has 1000 rows and the `departments` table has 9 rows. This is because there are 3 employees in the `employees` table having `NULL` value in the `dept_no` column.

You can't use `dept_no` alone in the `SELECT` clause, because it's ambiguous (both tables have `dept_no` column), so we need to specify the table name as in `e.dept_no`.

Another example if we want to get the grade of each student in each course, we will work on 3 tables (`students`, `courses`, and `students_course`) and use the primary and foreign keys to join the tables.

```
1 SELECT
2     student_name,
3     course_name,
4     grade
```

```
5 FROM
6   students s,
7   courses c,
8   students_course sc
9 WHERE
10  sc.student_no = s.student_no
11  AND c.course_no = sc.course_no;
```

To select which employee worked in which department, we can use the following query.

```
1 SELECT
2   first_name,
3   last_name,
4   dept_name
5 FROM
6   employees e,
7   departments d,
8   dept_emp de
9 WHERE
10  e.emp_no = de.emp_no
11  AND d.dept_no = de.dept_no
12 ORDER BY
13  e.emp_no;
```

2 Aggregate Functions

Aggregate functions work on a set of values and return a single value.

We will explain COUNT, SUM, AVG, MIN, and MAX functions.

To get the minimum, maximum, average, and total salary of employees, we can use the following query.

```
1 SELECT
2   MIN(salary),  -- Minimum salary between all employees
3   MAX(salary),  -- Maximum salary between all employees
4   SUM(salary),  -- Total salary for all employees
5   AVG(salary),  -- Average salary for all employees
6   COUNT(salary) -- Number of employees
7 FROM
8   employees;
```

This query will return 5 columns and 1 row. The query above will operate on the whole table, since there is no selected column in the SELECT clause.

2.1 GROUP BY & HAVING

To select another column next to the aggregate functions, we need to group the rows based on that column using GROUP BY.

So for example if we want to get the total salary for each gender we need to GROUP BY the gender column.

```
1 SELECT
2     gender,
3     SUM(salary)
4 FROM
5     employees
6 GROUP BY
7     gender;
```

Note 1:



Any column selected next to an aggregate function should be included in the GROUP BY clause.

To get the sum of salaries for each department, we can use the following query.

```
1 SELECT
2     dept_name,
3     SUM(salary)
4 FROM
5     employees e,
6     departments d
7 WHERE
8     e.dept_no = d.dept_no
9 GROUP BY
10    dept_name;
```

To get the sum of salary and count of employees in each department:

```
1 SELECT
2     dept_name,
3     SUM(salary),
4     COUNT(emp_no)
5 FROM
6     employees e,
7     departments d
8 WHERE
9     e.dept_no = d.dept_no
10 GROUP BY
11    dept_name;
```

Note 2:



When applying a condition on an aggregate function put it inside HAVING clause not WHERE

To get the sum of salary and count of employees in each department where the total salary is more than 15,000,000:

```
1 SELECT
2     dept_name,
3     SUM(salary),
4     COUNT(emp_no)
5 FROM
6     employees e,
7     departments d
8 WHERE
9     e.dept_no = d.dept_no
10 GROUP BY
11     dept_name
12 HAVING
13     SUM(salary) > 15000000
14 ORDER BY
15     SUM(salary) DESC;
```

To get the number of courses for each department:

```
1 SELECT
2     student_name,
3     COUNT(course_no)
4 FROM
5     students_course sc,
6     students s
7 WHERE
8     s.student_no = sc.student_no
9 GROUP BY
10     student_name;
```

To get students with age > 25, we used WHERE not HAVING because *we are filtering on a column not on a function*:

```
1 SELECT
2     student_name,
3     COUNT(course_no)
4 FROM
5     students_course sc,
6     students s
7 WHERE
8     s.student_no = sc.student_no
9     AND age > 25
10 GROUP BY
11     student_name;
```

In the query below we are grouping by `student_no` not `student_name` because we may have two students with the same name which will cause them to be grouped together as one student:

```
1 SELECT
2     student_name,
3     SUM(grade),
4     COUNT(course_no)
5 FROM
6     students s,
```

```
7      students_course sc
8 WHERE
9      s.student_no = sc.student_no
10 GROUP BY
11      s.student_no,
12      student_name;
```

To get the number of salaries each employee has in the `salaries` table, we can use the following query.

```
1 SELECT
2     e.emp_no,
3     first_name,
4     last_name,
5     COUNT(s.salary) AS emp_salary_count
6 FROM
7     employees e,
8     salaries s
9 WHERE
10    e.emp_no = s.emp_no
11 GROUP BY
12    e.emp_no,
13    first_name,
14    last_name
15 ORDER BY
16    emp_no;
```

And if we want to get employees with more than 10 different salary records, we can add `HAVING`:

```
1 SELECT
2     e.emp_no,
3     first_name,
4     last_name,
5     COUNT(s.salary) AS emp_salary_count
6 FROM
7     employees e,
8     salaries s
9 WHERE
10    e.emp_no = s.emp_no
11 GROUP BY
12    e.emp_no,
13    first_name,
14    last_name
15 HAVING
16    COUNT(s.salary) > 10
17 ORDER BY
18    emp_no;
```

3 UNION ALL & UNION

1. `UNION ALL` is used to combine the results of two or more `SELECT` statements, and it returns all rows from the combined queries.

2. UNION is used to combine the results of two or more SELECT statements, and it returns only *distinct* rows from the combined queries. For example if a row is returned by both queries, it will be shown only once.

Note 3:

For the union to work we must satisfy the following conditions:

- The number of columns in each query must be the same
- The data type of each column must be the same

You may have two different columns but with the same data type the union will work for that case but the data will be meaningless, so you should be careful when using UNION

If we have those two queries:

```
1 SELECT
2   first_name,
3   last_name,
4   salary
5 FROM
6   employees
7 WHERE
8   salary > 105000;
```

```
1 SELECT
2   first_name,
3   last_name,
4   salary
5 FROM
6   employees
7 WHERE
8   salary < 120000;
```

Each one will return a different table and we want to combine the results of both queries, we can use UNION ALL:

```
1 SELECT
2   first_name,
3   last_name,
4   salary
5 FROM
6   employees
7 WHERE
8   salary > 105000
9 UNION ALL
10 SELECT
11   first_name,
12   last_name,
13   salary
14 FROM
15   employees
16 WHERE
17   salary < 120000;
```

We can get the result of the query above using WHERE ... OR ...:

```
1 SELECT
2   first_name,
```

```
3  last_name,  
4  salary  
5  FROM  
6  employees  
7  WHERE  
8  salary > 105000  
9  OR salary < 130000  
10 ORDER BY  
11 first_name,  
12 last_name;
```

4 INTERSECT & EXCEPT

1. INTERSECT returns only the rows that are returned by both queries.
2. EXCEPT returns only the rows that are returned by the first query and not by the second query.

For example if we have those two queries:

```
1  SELECT  
2  first_name,  
3  last_name,  
4  salary  
5  FROM  
6  employees  
7  WHERE  
8  salary > 105000;
```

```
1  SELECT  
2  first_name,  
3  last_name,  
4  salary  
5  FROM  
6  employees  
7  WHERE  
8  salary < 130000  
9  ORDER BY  
10 first_name,  
11 last_name;
```

We can use INTERSECT to get the employees who have a salary between 105,000 and 130,000:

```
1  SELECT  
2  first_name,  
3  last_name,  
4  salary  
5  FROM  
6  employees  
7  WHERE  
8  salary > 105000  
9  INTERSECT  
10 SELECT  
11 first_name,  
12 last_name,  
13 salary  
14 FROM  
15 employees  
16 WHERE  
17 salary < 130000
```



```
18 ORDER BY
19     first_name,
20     last_name;
```

We can also use BETWEEN to get the same result:

```
1 SELECT
2     first_name,
3     last_name,
4     salary
5 FROM
6     employees
7 WHERE
8     salary BETWEEN 105000 AND 130000
9 ORDER BY
10     first_name,
11     last_name;
```

If we have a number of departments and we want to get departments that have no employees, we can use EXCEPT:

```
1 SELECT
2     dept_no,
3     dept_name
4 FROM
5     departments;
```

The query above will return all departments, and we want to get the departments that have no employees, so we will select unique `dept_no` from the `employees` table and use EXCEPT to get the departments that have no employees:

```
1 SELECT
2     dept_no,
3     dept_name
4 FROM
5     departments
6 EXCEPT
7 SELECT DISTINCT
8     d.dept_no,
9     dept_name
10 FROM
11     employees e,
12     departments d
13 WHERE
14     e.dept_no = d.dept_no;
```

Here the query after EXCEPT will return the departments that have employees, then we use EXCEPT to subtract those departments from all departments.